

# 08-Concurrency-Monitors-Multiprocessing

November 21, 2019

## 1 Monitors

A monitor is an object that only one thread can execute methods of this object at a time.

In Java Threads there is a synchronized modifier making a method a monitor like mutually exclusive.

In python you can use mutexes in the class to achieve that.

```
In [1]: from threading import Lock, Thread, Condition, RLock, Semaphore
```

```
class Monitor:
    def __init__(self):
        self.mutex = RLock()
        self.count = 0

    def increment(self):
        with self.mutex:
            self.count += 1

    def get(self):
        with self.mutex:
            return self.count
```

```
In [3]: def f(mon):
        for i in range(10000):
            mon.increment()

m = Monitor()

t1 = Thread(target=f, args=(m,))
t2 = Thread(target=f, args=(m,))
t3 = Thread(target=f, args=(m,))
t4 = Thread(target=f, args=(m,))

t1.start()
```

```

t2.start()
t3.start()
t4.start()
t1.join()
t2.join()
t3.join()
t4.join()
print(m.get())

```

40000

```

In [1]: class Queue:
        def __init__(self, capacity):
            self.capacity = capacity
            self.queue = []
            self.mutex = RLock()
            self.notempty = Lock(0)

        def empty(self):
            return len(self.queue) == 0

        def full(self):
            return len(self.queue) == self.capacity

        def enqueue(self, val):
            with self.mutex:
                while self.full():
                    time.sleep(1)
                self.queue.append(val)
                if len(self.queue) == 1:
                    self.notempty.release()

        def dequeue(self):
            with self.mutex:
                while self.empty():
                    self.mutex.release()
                    self.notempty.acquire()
                    self.mutex.acquire()
                a=self.queue[0]
                del self.queue[0]

```

## 2 Producer Consumer

- A queue which is accessed by two (or more) threads. One end a producer thread inserts items, the other end, consumer thread removes and processes the items.
- They work in an infinite loop.
- If queue is empty or full?

- In full and empty cases, they need to check it until queue has an empty slot or has an item respectively. Polling in an infinite loop wastes too much CPU
- Busy waiting is not a good idea:

```
while queue.empty():
    time.sleep(1)  # response time will be slow
    pass
```

- Use synchronization methods semaphores or similar to make other end know that queue is ready (not full or not empty)

### 3 Condition Variables

- In a monitor, condition variables let threads to signal each other while keeping the monitor semantics (only one thread inside).

```
c = Condition(mutex)
c.wait()
```

wait does:

```
c.mutex.release()
# block on condition
# when unblocked:
c.mutex.acquire()
```

Typical usage:

```
c.acquire() # or acquire mutex of c on construction
.....
while actual condition:
    c.wait()
```

The notifier cannot guarantee that the condition holds semantically and notified thread can directly assume condition holds.

`c.notify()` will unblock one of the threads blocking on condition.

`c.notifyAll()` will unblock all of them. However they still wait on the mutex after unblocking. They enter monitor one at a time. asdasd

```
In [7]: import random
import time

class PCQueue:
    def __init__(self, capacity=10):
        self.mutex=RLock()
        self.queue = []
        self.capacity = capacity
        self.notempty = Condition(self.mutex)
        self.notfull = Condition(self.mutex)
```

```

def empty(self):
    with self.mutex:
        return len(self.queue) == 0
def full(self):
    with self.mutex:
        return len(self.queue) == self.capacity
def enqueue(self, item):
    with self.mutex:
        while len(self.queue) == self.capacity:
            print("queue is full, waiting")
            self.notfull.wait()

        self.queue.append(item)
        self.notempty.notify()

def dequeue(self):
    with self.mutex:
        while len(self.queue) == 0:
            print("queue is empty, waiting")
            self.notempty.wait()

        val = self.queue[0]
        del self.queue[0]
        self.notfull.notify()
        return val

def producer(pcq):
    for i in range(30):
        time.sleep(0.15+random.random()*0.15)
        pcq.enqueue(random.randint(0,100))
        print("enqueued")
    print("producer finished")

def consumer(pcq):
    for i in range(30):
        time.sleep(0.05+random.random()*0.2)
        print("dequeued ", pcq.dequeue())
    print("consumer finished")

q = PCQueue()

prod = Thread(target=producer, args=(q,))
cons = Thread(target=consumer, args=(q,))
prod.start()
cons.start()
prod.join()
cons.join()

```

queue is empty, waiting  
enqueued  
dequeued 72  
queue is empty, waiting  
enqueueddequeued 4

enqueued  
dequeued 94  
queue is empty, waiting  
enqueueddequeued

11  
queue is empty, waiting  
enqueued

dequeued 12  
queue is empty, waiting  
enqueued

dequeued 17  
queue is empty, waiting  
enqueued

dequeued 76  
enqueued  
dequeued 96

enqueued  
dequeued 87  
enqueued

dequeued 81  
queue is empty, waiting  
enqueued

dequeued 40  
queue is empty, waiting  
enqueued

dequeued 0  
enqueued  
dequeued 30

queue is empty, waiting  
enqueueddequeued  
98

enqueueddequeued  
28  
queue is empty, waiting

enqueued  
dequeued 26  
queue is empty, waiting

enqueued  
dequeued 73  
queue is empty, waiting

enqueued  
dequeued 64

```

queue is empty, waiting
enqueueddequeued
36
queue is empty, waiting
dequeued enqueued 71

queue is empty, waiting
enqueued
dequeued 39
queue is empty, waiting
enqueueddequeued
37
queue is empty, waiting
enqueueddequeued
86
queue is empty, waiting
enqueued
dequeued 57
queue is empty, waiting
enqueued
dequeued 5
enqueued
dequeued 85
queue is empty, waiting
enqueued
dequeued 40
queue is empty, waiting
enqueued
dequeued 29
queue is empty, waiting
enqueued
dequeued 24
enqueued
producer finished
dequeued 44
consumer finished

```

```

In [8]: prod = Thread(target=producer, args=(q,))
        prod2 = Thread(target=producer, args=(q,))
        cons = Thread(target=consumer, args=(q,))
        cons2 = Thread(target=consumer, args=(q,))
        prod.start()
        cons.start()
        prod2.start()
        cons2.start()

        prod.join()

```

```

        cons.join()
        prod2.join()
        cons2.join()

queue is empty, waiting
enqueued
dequeued 60
queue is empty, waiting
enqueued
dequeued 70
queue is empty, waiting
queue is empty, waiting
enqueued
dequeued 18
enqueued
dequeued 31
enqueued
dequeued 94
enqueued
enqueued
dequeued 38
dequeued 85
queue is empty, waiting
queue is empty, waiting
enqueued
dequeued 49
enqueueddequeued
18
enqueued
dequeued 71
queue is empty, waiting
enqueueddequeued
69
enqueued
dequeued 47
queue is empty, waiting
queue is empty, waiting
enqueued
dequeued 94
dequeued enqueued 43

queue is empty, waiting
enqueueddequeued enqueueddequeued

9341

queue is empty, waiting
enqueued

```

dequeued 61  
queue is empty, waiting  
enqueueddequeued  
19  
queue is empty, waiting  
queue is empty, waiting  
enqueueddequeued  
82  
enqueueddequeued  
5  
queue is empty, waiting  
queue is empty, waiting  
dequeued enqueued  
6  
enqueued  
dequeued 45  
queue is empty, waiting  
queue is empty, waiting  
enqueueddequeued  
92  
enqueueddequeued  
97  
queue is empty, waiting  
enqueueddequeued queue is empty, waitingenqueued  
8

dequeued 40  
queue is empty, waiting  
queue is empty, waiting  
enqueueddequeued  
88  
enqueued  
dequeued 82  
queue is empty, waiting  
dequeued enqueued  
96  
enqueued  
dequeued 100  
enqueued  
dequeued 7  
queue is empty, waiting  
enqueueddequeued  
64  
queue is empty, waiting  
enqueued  
dequeued 91  
queue is empty, waiting



enqueued  
dequeued 29  
queue is empty, waiting  
enqueueddequeued 71

enqueued  
dequeued 22  
queue is empty, waiting  
enqueueddequeued 64

queue is empty, waiting  
enqueued  
dequeued 81  
queue is empty, waiting  
enqueued  
dequeued 84  
queue is empty, waiting  
enqueued  
dequeued 3  
enqueued  
dequeued 44  
queue is empty, waiting  
queue is empty, waiting  
dequeued enqueued  
30  
queue is empty, waitingenqueueddequeued

59  
queue is empty, waiting  
enqueued  
dequeued 89  
enqueueddequeued

8  
queue is empty, waiting  
enqueued  
dequeued 27  
queue is empty, waiting  
queue is empty, waiting  
enqueueddequeued

21  
enqueueddequeued queue is empty, waiting  
2

queue is empty, waiting  
enqueued  
dequeued 53  
enqueueddequeued 46

```

queue is empty, waiting
enqueueddequeued queue is empty, waiting

99
enqueueddequeued
23
queue is empty, waiting
enqueued
dequeued 90
queue is empty, waiting
enqueueddequeued
56
queue is empty, waiting
queue is empty, waiting
enqueued
dequeued 88
enqueueddequeued
42
queue is empty, waiting
enqueued
dequeued 14
enqueueddequeued queue is empty, waiting

74
consumer finished
enqueued
dequeued producer finished
43
enqueued
producer finished
dequeued 48
consumer finished

```

## 4 multiprocessing.Queue and Queue modules

```

from multiprocessing import ..., Queue
from threading import ... import Queue
A synchronized object in shared memory. All blocking/unblocking is already implemented.

```

```

In [9]: from multiprocessing import Queue, Process

```

```

q = Queue(10)

def producer(q):
    for i in range(100):
        q.put(i)

```

```
def consumer(q):
    for i in range(100):
        item = q.get()
        print(item)

prod=Process(target=producer, args=(q,))
cons=Process(target=consumer, args=(q,))
prod.start()
cons.start()
prod.join()
cons.join()
```

0  
1  
2  
3  
4  
5  
6  
7  
8  
9  
10  
11  
12  
13  
14  
15  
16  
17  
18  
19  
20  
21  
22  
23  
24  
25  
26  
27  
28  
29  
30  
31  
32  
33  
34  
35

36  
37  
38  
39  
40  
41  
42  
43  
44  
45  
46  
47  
48  
49  
50  
51  
52  
53  
54  
55  
56  
57  
58  
59  
60  
61  
62  
63  
64  
65  
66  
67  
68  
69  
70  
71  
72  
73  
74  
75  
76  
77  
78  
79  
80  
81  
82  
83

84  
85  
86  
87  
88  
89  
90  
91  
92  
93  
94  
95  
96  
97  
98  
99

## 5 Process/Thread Pools

- $[i_1, \dots, i_N]$  items and apply  $f()$  to all of them in parallel to get  $[f(i_1), f(i_2), \dots, f(N)]$  as a result.
- creating  $N$  threads/process looks logical but resources are limited.  $N == 4$  it is ok but if  $N == 10000$ ?
- Instead create  $M$  processes and compute in groups of  $M$ .

```
In [11]: from multiprocessing import Pool
```

```
    pool = Pool(8)

    def f(i):
        time.sleep(0.2+0.3*random.random())
        return i*i

    g = pool.map(f, [i for i in range(100)])
    print(g)
```

[0, 1, 4, 9, 16, 25, 36, 49, 64, 81, 100, 121, 144, 169, 196, 225, 256, 289, 324, 361, 400, 441,

Implement your own pool?

## 6 Deadlock and Dining Philosophers

```
In [2]: from threading import *
        from time import *
        from random import *
```

```

STARTED = 0
THINKING = 1
HUNGRY = 2
EATING = 3
EXITTED = 4

stmess = "0?-*X"

class Philosopher(Thread):
    def __init__(self, id, forks, states, updated):
        Thread.__init__(self)
        self.id = id
        self.left = forks[0]
        self.right = forks[1]
        self.states = states
        self.states[id] = STARTED
        self.term = False
        self.updated = updated
    def terminate(self):
        self.term = True
    def run(self):
        for i in range(10):
            if self.term:
                break
            # print self.id, " is thinking"
            with self.updated:
                self.states[self.id] = THINKING
                self.updated.notify()

            sleep(random()*1)
            with self.updated:
                self.states[self.id] = HUNGRY
                self.updated.notify()
            if self.id % 2 == 0:
                self.left.acquire()
                self.right.acquire()
            else:
                self.right.acquire()
                self.left.acquire()
            with self.updated:
                self.states[self.id] = EATING
                self.updated.notify()
            # print self.id, " is eating"
            sleep(random()*4)
            self.left.release()
            self.right.release()
            with self.updated:

```

```

        self.states[self.id] = EXITTED

print("Enter number of philosopher: ", end='')
n = int(input())

forks = [Lock() for i in range(n)]

phils = []

states = [0 for i in range(n)]

updated = Condition()

for i in range(n):
    phils.append( Philosopher(i,(forks[i],forks[(i+1)%n]),states, updated) )

for phil in phils:
    phil.start()

while True:
    eflag = True
    for i in range(n):
        if states[i] != EXITTED:
            eflag = False
            print(stmess[states[i]],end='')
    print()
    if eflag:
        break
    try:
        with updated:
            updated.wait()
    except KeyboardInterrupt:
        for phil in phils:
            phil.terminate()

for phil in phils:
    phil.join()

```

Enter number of philosopher: 4

-----

RuntimeError

Traceback (most recent call last)

<ipython-input-2-7f73b699911b> in <module>()

```

68
69 for phil in phils:
---> 70     phil.start()
71
72 while True:

/usr/lib/python3.5/threading.py in start(self)
842         _limbo[self] = self
843     try:
--> 844         _start_new_thread(self._bootstrap, ())
845     except Exception:
846         with _active_limbo_lock:

```

```

RuntimeError: can't start new thread

```

## 7 Synchronizing/Watching a Thread/Process

- Have a condition variable for synchronization.
- Send it to Thread/Process
- In the watcher wait for it
- When the model/state changes in thread/process, notify the condition variable.

Call a function in a thread asynchronously (assume there are multiple threads, join() only joins one of them):

```

def f(x):
    return x*x

def call(c):
    with c[3]:
        c[0] = c[1](c[2])
        c[3].notify()

c = Condition()
# (result, function, input, condition)
result=[None, f, 15, c]
t = Thread(target=call, args=(result))
with c:
    c.wait()

In [12]: from threading import Thread,Condition,Lock
import time

class AsyncCall(Thread):
    def __init__(self,func,args):

```



```

        super().__init__()
        self.func = func
        self.args=args
        self.cond=Condition()
        self.ready = False
        self.start()
    def run(self):
        self.value = self.func(self.args)
        with self.cond:
            self.ready = True
            self.cond.notifyAll()
    def wait(self):
        with self.cond:
            while not self.ready:
                self.cond.wait()

def f(x):
    time.sleep(3)
    return x*x

c = AsyncCall(f,10000)
print("I can do usefull stuff here...")
c.wait()
print(c.value)

```

```

I can do usefull stuff here...
100000000

```

## 8 Concurrency Overview

- Watch race conditions! Use locks/semaphores to protect them
- Watch deadlocks. Be careful when holding a lock and try to acquire another.
- Never make assumptions about timing!. Timing of a thread becoming ready, calling some heavy function. OS/PL scheduler can behave undeterministically.
- Never busy wait!
- Use monitors and condition variables when you need higher level abstractions of synchronization. a monitor queue for producer consumer
- Be careful about if your data is shared or not! multiprocessing: not shared, use Value/Array/Queue threading: all globals and **object** parameters are shared
- Be careful about Global Interpreter Lock: If task is I/O intensive threading should work. but if it has cpu intensive mostly -> no parallelism. threading: lightweight, shared variables default, easy to manage but worse parallelism multiprocessing: parallel, but more expensive, needs explicit shared variables

- If a process/thread has behavior, implement as a derived class

```
class myclass(Process):    or class myclass(Thread)
```

call `super().__init__()` in constructor override `run()` method. if a simple function, just start it.

- If only synchronization is required, your classed can be anything, implement a monitor